

Meridian Supply Co.

demo-store.example · GMV band \$1M–\$5M · 2026-06-15

78 / 100

AGENT READINESS SCORE

AGENT ACCESS (ROBOTS.TXT)		7.7/10
STRUCTURED DATA (JSON-LD)		13.9/15
PRODUCT FEEDS (PRODUCTS.JSON, SITEMAP)		10/10
LLMS.TXT		0/5
CART REACHABLE		20/20
CHECKOUT REACHABLE		20/20
LIVE AGENT OUTCOMES		6.7/20

EXECUTIVE SUMMARY

NOTE Agent-attributed orders represent \$72,396–\$112,801 of annualized revenue (measured from your own orders).

HIGH With the checkout blockers found here, an estimated \$36,198–\$56,401/yr of agent demand is at risk as the channel grows (scenario, not a measurement).

HIGH Worst qualifying month for disputes was 2026-04: 0.83% (13/1562), above Visa's monitoring thresholds.

HIGH 3 of 13 agent user-agents are blocked by robots.txt (GPTBot, ClaudeBot, Bytespider).

MEDIUM Agent-attributed orders dispute at 3.06% vs 0.41% for human orders.

NOTE Agent order share rose from 6.1% (2026-03) to 11.1% (2026-06) over the window.

NOTE The automated probe reached the checkout information page.

NOTE No llms.txt is published — an emerging convention that guides agents through the catalog.

MEDIUM 2 of 8 sampled product pages have incomplete Product/Offer fields (offer missing priceCurrency; offer missing availability; missing canonical; product missing sku/gtin).

MEDIUM Perplexity abandoned the purchase at the variant step.

MEDIUM Claude abandoned the purchase at the checkout step.

HOW YOU COMPARE

Your discovery-readiness score of **79/100** — the public-surface signals every store exposes (robots, structured data, feeds, llms.txt) — places you in the **bottom 48%** of a cohort of 50 stores (median 85.5). 26 score higher.

Cohort: Top 50 Shopify-powered stores, June 2026. These are large, well-known stores, so a low rank here is a conservative read — the typical store scores lower. Discovery (public-surface) only; transaction-layer and dispute benchmarks are not cohort-derived.

WHAT AGENT COMMERCE IS WORTH TO YOU

AGENT-ATTRIBUTED REVENUE (ANNUALIZED)

\$72,396–

\$112,801

measured: \$17,851–\$27,814 over
90d

DISPUTE HANDLING COST (ANNUALIZED)

\$1,947–\$3,244

+ \$6,343 disputed order value

FORWARD EXPOSURE (SCENARIO)

\$36,198–

\$56,401/yr

if the checkout blockers below
persist as agent traffic grows

ASSUMPTIONS

- Annualized figures multiply the 90-day window by 4.06 (365/90); they assume the window is representative.
- Captured agent revenue is the store's OWN measured GMV for agent-attributed orders. Floor = confirmed + high-confidence tiers; ceiling adds the heuristic tier, which may include datacenter/VPN traffic that is not a true agent.
- Dispute cost uses an assumed per-dispute handling fee of \$15–\$25 (industry-typical; not measured) plus the merchant-reported disputed amount, which covered 91% of disputes.
- Forward exposure is a SCENARIO, not a measurement: 2 of 4 tested agent paths failed to reach checkout (rate 50%). It estimates demand at risk as the agent channel grows; it is NOT a claim that already-captured revenue is lost.

DISPUTE EXPOSURE (VISA VAMP)

Worst qualifying month (2026-04): 13 disputes / 1562 orders = 0.83% — above Visa's standard threshold.

VAMP is assessed monthly. We show every month but suppress the ratio for any month under 200 orders so a low-volume month can't manufacture an alarming number. Denominator is total orders (a proxy that overstates the true ratio: it includes non-Visa tender and inquiries).

MONTH	ORDERS	DISPUTES	RATIO	BAND
2026-03	799	3	0.38%	ok
2026-04	1,562	13	0.83%	above std
2026-05	1,905	12	0.63%	above std
2026-06	814	4	0.49%	ok

Disputed order value over the window: **\$1,564** (covering 91% of disputes that carried an amount). Of that, \$366 is agent-attributed. Annualized dispute handling cost at an assumed fee band: **\$1,947–\$3,244/yr.**

Agent-attributed orders dispute at **3.06%** vs **0.41%** for human orders (13 agent-side disputes across 425 agent orders).

Setting aside the 13 agent-attributed disputes, the blended 90-day ratio falls from **0.63%** (above standard) to **0.37%** (ok) — an arithmetic illustration of the agent channel's weight, not a causal claim.

BY SEGMENT

CLASS	ORDERS	DISPUTES	DISPUTE RATE
Confirmed channel	33	0	n low
High-confidence agent	238	5	2.10%
Heuristic agent	154	8	5.19%
Human	4,655	19	0.41%

CAN EACH AGENT BUY FROM YOU?

AGENT	DISCOVER	READ PRODUCT	REACH CART	REACH CHECKOUT	LIVE OUTCOME
ChatGPT (OpenAI)	yes	yes	yes	yes	completed
Perplexity	yes	yes	NO	NO	abandoned at variant
Claude (Anthropic)	yes	yes	yes	NO	abandoned at checkout
Google (Gemini/AI) INFERRED	yes	yes	partial	partial	not live-tested
Amazon (Rufus) INFERRED	yes	yes	partial	partial	not live-tested
Apple Intelligence INFERRED	yes	yes	partial	partial	not live-tested
Meta AI INFERRED	yes	yes	partial	partial	not live-tested

Rows tagged **INFERRED** were not live-tested (only ChatGPT, Perplexity and Claude are run as live purchases today). Their cart/checkout cells are inferred from robots access and the shared automated probe, not a per-agent purchase attempt. "Discover" = allowed by robots.txt; "Read product" = required structured-data fields present.

AGENT-SHARE TRAJECTORY

Agent order share moved from **6.1%** (2026-03) to **11.1%** (2026-06).




MONTH	ORDERS	AGENT ORDERS	AGENT GMV SHARE
2026-03	799	49 · 6.1%	6.4%
2026-04	1562	111 · 7.1%	7.7%
2026-05	1905	175 · 9.2%	9.8%
2026-06	814	90 · 11.1%	10.8%

A 90-day window is a short baseline; treat the slope as directional. Months with low order counts are noisy.

ORDER CLASSIFICATION

5,080 orders (\$315,985 GMV) over the last 90 days.

Agent-attributed order share: **5.3%** (floor, strong signals) to **8.4%** (ceiling, heuristic-inclusive). GMV share: 5.6%–8.8%.



CLASS	ORDERS	SHARE	GMV	AOV	REFUND RATE
Confirmed channel	33	0.6%	\$1,563	\$47	2 (n low)
High-confidence agent	238	4.7%	\$16,288	\$68	7.1%
Heuristic agent	154	3.0%	\$9,963	\$65	8.4%
Human	4,655	91.6%	\$288,171	\$62	4.7%

DISTINCT ORDER SOURCES

SOURCE_NAME	APP_ID	ORDERS
web	580111	5,047
chatgpt ASSISTANT	580111	33

Refund rate counts refunded/partially-refunded/voided orders (a leading indicator that surfaces faster than formal disputes). Rates are suppressed where a class has fewer than 100 orders.

DISCOVERY LAYER

10/13 agent user-agents allowed by robots.txt. Blocked: `GPTBot` (Disallow: /), `ClaudeBot` (Disallow: /), `Bytespider` (Disallow: /).

FEED	STATUS	NOTES
<code>products.json</code>	pass	products.json accessible
<code>sitemap.xml</code>	pass	sitemap ok, 8 product urls sampled
<code>llms.txt</code>	fail	llms.txt returned 404

PRODUCT PAGE	JSON-LD FIELDS	PROBLEMS
<code>/products/field-jacket</code>	5/5	clean
<code>/products/canvas-tote</code>	5/5	clean
<code>/products/trail-bottle</code>	5/5	clean
<code>/products/wool-beanie</code>	5/5	clean
<code>/products/camp-mug</code>	5/5	clean
<code>/products/dry-bag</code>	5/5	clean
<code>/products/trek-socks</code>	3/5	offer missing priceCurrency; offer missing availability; missing canonical
<code>/products/gift-card</code>	4/5	product missing sku/gtin; missing og:image

TRANSACTION LAYER

AGENT	STAGE REACHED	BLOCKER / NOTES
Automated probe	checkout (info page) (41s)	popup (product_page)
ChatGPT	completed (stopped pre-payment)	Smooth run; agent handled the variant picker without help.
Perplexity	abandoned at variant	Variant radio group has no accessible labels; agent could not pick a size.
Claude	abandoned at checkout	Stacked modals; second one has no close button visible to the agent.

WHERE AGENTS DROP OFF

The largest cluster of failures is at **Discovery (robots / feeds)** — fix #1 below targets this stage.

Discovery (robots / feeds)		3
Product page (read details)		3
Variant selection		1
Reach checkout		1

- **Discovery (robots / feeds):** 3 agent UAs blocked by robots.txt
- **Product page (read details):** 2 product pages with structured-data gaps; probe: popup
- **Variant selection:** perplexity abandoned
- **Reach checkout:** claude abandoned

FIX LIST

#	FIX (RANKED BY IMPACT ÷ EFFORT)	IMPACT	EFFORT
1	Allow agent user-agents in robots.txt (3 currently blocked) Blocked agents never see the catalog; one file edit restores discovery for every assistant.	5/5	1/5
2	Defer popups off product and cart pages Interstitials cover the buy controls; agents rarely recover after a popup mis-close.	4/5	2/5
3	Server-render complete JSON-LD Product/Offer fields Price, currency, availability and identifiers are what agents quote before visiting — and they must be in the initial HTML, since most crawler-class agents don't run JavaScript.	4/5	2/5
4	Add canonical and OpenGraph tags to product pages Cheap wins for deduplication and previews in assistant answers.	2/5	1/5
5	Publish an llms.txt index Emerging convention; cheap to add and increasingly read by assistants.	2/5	1/5
6	Make the variant picker machine-readable (labels + native inputs) Agents abandoned at variant selection; accessible labels fix agents and screen readers alike.	4/5	3/5

COPY-PASTE FIXES

ROBOTS.TXT — UNBLOCK 3 AGENT USER-AGENTS

Currently blocked: GPTBot, ClaudeBot, Bytespider. Add these stanzas to your robots.txt (Shopify: Online Store > themes > edit code > robots.txt.liquid).

```
# Allow AI shopping agents to read your storefront
User-agent: GPTBot
User-agent: ClaudeBot
User-agent: Bytespider
Allow: /

# (keep your existing rules for other crawlers below)
```

SERVER-RENDERED JSON-LD PRODUCT/OFFER

Serve this in the initial HTML (server-rendered), not injected by JavaScript — most crawler-class agents don't run JS. Affected pages: /products/trek-socks, /products/gift-card.

```
<script type="application/ld+json">
{
  "@context": "https://schema.org",
  "@type": "Product",
  "name": "<product title>",
  "image": ["https://demo-store.example/<product-image>.jpg"],
  "sku": "<your SKU or GTIN>",
  "brand": { "@type": "Brand", "name": "Meridian Supply Co." },
  "offers": {
    "@type": "Offer",
    "price": "<numeric price, no currency symbol>",
    "priceCurrency": "USD",
    "availability": "https://schema.org/InStock",
    "url": "https://demo-store.example/products/<handle>"
  }
}
</script>
```

LLMS.TXT — A GUIDED INDEX FOR AI AGENTS

Serve this at <https://demo-store.example/llms.txt>. Emerging convention; cheap to add and increasingly read by assistants.

```
# Meridian Supply Co.
> Meridian Supply Co. storefront. Products, prices, and policies for AI shopping
agents.

## Products
- Catalog (JSON): https://demo-store.example/products.json
- Sitemap: https://demo-store.example/sitemap.xml

## Policies
- Shipping: https://demo-store.example/policies/shipping-policy
- Returns: https://demo-store.example/policies/refund-policy

## Contact
- owner@demo-store.example
```

APPENDIX — HOW WE SCORE

The **Agent Readiness Score** runs 0–100, split into a **Discovery half (40 pts)** — can agents find and read your catalog — and a **Transaction half (60 pts)** — can they actually buy. Each component is earned proportionally (not pass/fail), then summed.

COMPONENT	MAX PTS	HOW POINTS ARE EARNED
Discovery · Agent access (robots.txt)	10	allowed agent user-agents ÷ 13, times the weight
Discovery · Structured data (JSON-LD)	15	mean across sampled product pages of (required Product/Offer fields present ÷ 5), times the weight
Discovery · Product feeds	10	half for a reachable /products.json, half for a valid sitemap.xml
Discovery · llms.txt	5	full if /llms.txt is published
Transaction · Cart reachable	20	full if the automated probe reached the cart
Transaction · Checkout reachable	20	full if it reached the checkout information page
Transaction · Live agent outcomes	20	successful live agent runs ÷ total live runs, times the weight

"Required Product/Offer fields" = price, priceCurrency, availability, SKU or GTIN, and image (5 fields), read from server-rendered JSON-LD.

DISCOVERY READINESS SCORE

The four Discovery rows rescaled to 0–100 (discovery points ÷ 40 × 100). It needs no order data, so it is the headline for readiness-only audits and the basis for the cross-store percentile — an apples-to-apples comparison whether or not a store shared its order data.

HOW WE CLASSIFY ORDERS

- **Confirmed channel** — a non-web `source_name / app_id` recorded by Shopify. Not inferred; zero ambiguity.
- **High-confidence agent** — a published agent user-agent, an assistant referrer (chatgpt.com, perplexity.ai, claude.ai, ...), or an assistant `utm_source`. Strong signals, but referrers get stripped, so this undercounts.
- **Heuristic agent** — headless-browser markers or datacenter IP ranges (AWS/GCP/Azure). Signals, not proof: VPN users can land here, so we treat this tier as the ceiling and never the headline.
- **Human** — the default. Agents that perfectly mimic a consumer browser are invisible in order data, so true agent share is a floor, likely higher than reported.

DISPUTE & VAMP MATH

Dispute ratio = disputes ÷ orders over the 90-day window, shown per month with minimum-volume suppression so a low-count month can't manufacture an alarming rate. Visa assesses VAMP monthly on settled VisaNet transactions via your acquirer, so treat our figure as a directional proxy, not the official number. Dollar figures use the merchant-reported disputed amount plus an explicitly-labeled, assumed per-dispute fee band — shown as a range, never a point estimate.

CAVEATS

- Discovery checks read the **initial HTML response** and do not execute JavaScript. Structured data injected client-side reads as missing here — which is also how agents that don't run JS (most crawler-class agents) see the page.
- Per-agent reachability is inferred from robots access and one shared automated probe, except for live runs (ChatGPT/Perplexity/Claude and any assistant we run by hand), which are real purchase attempts. Inferred rows are labeled.
- Automated and live checks stop at the checkout information / payment step; no purchase is ever completed and no payment details are entered. CAPTCHA presence is recorded, never solved or bypassed.